

Jaswanthi B

Forward Deployed Engineer

San Jose, CA (Open to Relocation / Remote) | 408-357-0210 | jaswanthi@jobmailoffer.com | [LinkedIn](#) | [GitHub](#) | [Portfolio](#)

SUMMARY

Forward Deployed and AI Engineer with 7 years of experience building multi-agent LLM architectures, computer vision pipelines, and distributed data infrastructure across enterprise and research settings. Hands-on expertise engineering production RAG systems, orchestrating multi-agent workflows with LangGraph, and serving ViT/SAM models via Triton Inference Server and FastAPI. Architected Allstate's claims-intelligence platform, automating ingestion and telemetry cross-referencing to reduce triage time from 3.5 days to under 10 minutes while improving early fraud detection by 30%. Strong background in cloud data pipelines, automated PII redaction and governance, real-time streaming, and CI/CD infrastructure.

SKILLS

Languages & Frameworks: Python, SQL, TypeScript, React, PySpark, FastAPI
AI, Multi-Agent & Machine Learning: LangGraph, LangChain, Multi-Agent Orchestration, Retrieval-Augmented Generation (RAG), Vision Transformers (ViT), Segment Anything Model (SAM), Pinecone, TruLens
Model Serving & Inference: Triton Inference Server, vLLM, TorchServe, RESTful Microservices, Model Retraining Pipelines
Data Engineering & Streaming: AWS (Glue, Lambda, S3, Redshift), Apache Spark, Apache Kafka, Apache Airflow, SSIS, PostgreSQL, MS SQL Server
DevOps, Infrastructure & Security: Docker, Terraform, AWS CloudFormation, Jenkins CI/CD, Git, PII Redaction & Data Governance, Role-Based Access Control (RBAC)

EXPERIENCE

Allstate | Forward Deployed Engineer|(IL - Remote) Jun 2025 - Present

- Architected an end-to-end claims-intelligence platform that reduced manual claim-triage processing time from 3.5 days to under 10 minutes by unifying automated ingestion, multi-agent telemetry cross-referencing, computer vision inference, and auto-routing workflows.
- Built a unified ingestion pipeline in FastAPI to process driver voice memos, crash images, and IoT vehicle telemetry with sub-2-second ingestion times.
- Enforced data governance and privacy protocols by implementing automated PII redaction and access controls across ingested audio voice memos and unstructured accident statements.
- Connected real-time vehicle telematics with accident statements through a LangGraph multi-agent workflow; cross-referenced telemetry (e.g., GPS impact speed vs. stated impact speed) to identify discrepancies, boosting early fraud detection rates by 30%.
- Deployed Vision Transformer (ViT) and Segment Anything (SAM) models as microservices via Triton Inference Server and FastAPI for vehicle-damage assessment, achieving sub-second inference and 89% classification accuracy.
- Developed a React-based claims review interface displaying model confidence scores and bounding-box visualizations, auto-routing high-confidence claims while seamlessly flagging edge cases for human adjuster review as part of the broader platform.
- Integrated TruLens model monitoring and adjuster feedback tracking directly into the platform, leveraging override patterns for continuous retraining workflows that improved damage-prediction precision over time.

University of Houston | Research Assistant - Data Science & AI | Houston, TX Jan 2023 - May 2025
(Continued post-graduation)

- Designed a LangGraph-based multi-agent workflow for author disambiguation, citation-network analysis, and co-authorship mapping across 50,000+ publication records, reducing manual validation effort.
- Applied NLP techniques, contextual embeddings, and ORCID/OpenAlex APIs to resolve duplicate researcher profiles and improve identity matching across publication datasets.
- Created a RAG-based research assistant with LangChain and Pinecone, enabling research staff to search publication metrics and grant trends through natural-language queries.
- Implemented data pipelines for batch embedding generation, document chunking, and hybrid vector search, supporting semantic retrieval across research datasets.
- Conducted training sessions on Python, SQL, and system design for students and research staff while maintaining Tableau dashboards for institutional performance reporting.

Capgemini | Cloud Data Engineer | Bengaluru, India Mar 2022 - Dec 2022
Client: Chevron Corporation | (Pre-Graduation Co-Op / Full-Time)

- Built real-time streaming data pipelines using AWS Glue, Apache Spark, and Kafka to ingest raw topdrive sensor telemetry and downhole drilling metrics into Amazon S3 for early detection of drill-bit fatigue, enabling field engineers to schedule proactive bit replacements and prevent costly downhole tool failures.
- Developed PySpark ETL workflows to validate, clean, and transform operational logs before loading them into Amazon Redshift, improving data quality and query performance.
- Managed workflow scheduling and production monitoring with Apache Airflow, AWS Lambda, and Amazon CloudWatch, helping reduce recurring pipeline issues and improve overall reliability.
- Automated cloud infrastructure provisioning with Terraform and AWS CloudFormation and integrated Jenkins CI/CD pipelines to streamline deployments.

Virtual Infotech Solution | Data Analyst | India Oct 2019 - Mar 2022
(Part-Time / Intern)

- Centralized reporting across departments by migrating fragmented spreadsheet workflows into a unified SQL Server and Power BI reporting environment.
- Automated data validation and ETL workflows using SSIS and Python scripts, significantly reducing manual reporting errors and department discrepancies.
- Developed Power BI dashboards for customer acquisition, revenue churn, and supply-chain KPIs, driving departmental adoption of self-service analytics

EDUCATION

Master of Science (M.S.) in Data Science Dec 2024
University of Houston | Houston, TX, USA
Bachelor of Technology (B.Tech.) in Computer Science and Engineering Jun 2022
Institute of Aeronautical Engineering | Hyderabad, India

CERTIFICATIONS

AWS Certified Solutions Architect - Associate | Amazon Web Services
Infosys Certified Software Programmer | Infosys